

# Database Systems

(CS 355 / CE 373)

Dr. Umer Tariq  
Assistant Professor,  
Dhanani School of Science & Engineering,  
Habib University

# Acknowledgements

- Many slides have been borrowed from the official lecture slides accompanying the textbook:

Database System Concepts, (2019), Seventh Edition,  
Avi Silberschatz, Henry F. Korth, S. Sudarshan  
McGraw-Hill, ISBN 9780078022159

The original lecture slides are available at:


<https://www.db-book.com/>

- Some of the slides have been borrowed from the lectures by Dr. Immanuel Trummer (Cornell University). Available at: ([www.itrummer.org](http://www.itrummer.org))

# Outline: Week 6

- SQL: An Overview
- Data Definition in SQL
- Basic Structure of SQL Queries
- Aggregate Functions

# DBMS-based Approach



- Data Model
  - A collection of conceptual tools for describing data, data relationships, data semantics, and consistency constraints.

# The Relational Model

- The relational model uses a collection of tables to represent both data and the relationships among those data.

The diagram illustrates the *instructor* relation as a table. A bracket on the left side of the table is labeled "Table/ Relation" and "instructor". A box highlights the header row, with an arrow pointing to it labeled "Column / Attribute" and "dept\_name". Another box highlights the row containing the tuple (83821, Brandt, Comp. Sci., 92000), with an arrow pointing to it labeled "Row / Tuple" and "(83821, Brandt, Comp. Sci., 92000)".

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 12121     | Wu          | Finance          | 90000         |
| 15151     | Mozart      | Music            | 40000         |
| 22222     | Einstein    | Physics          | 95000         |
| 32343     | El Said     | History          | 60000         |
| 33456     | Gold        | Physics          | 87000         |
| 45565     | Katz        | Comp. Sci.       | 75000         |
| 58583     | Califieri   | History          | 62000         |
| 76543     | Singh       | Finance          | 80000         |
| 76766     | Crick       | Biology          | 72000         |
| 83821     | Brandt      | Comp. Sci.       | 92000         |
| 98345     | Kim         | Elec. Eng.       | 80000         |

Figure 2.1 The *instructor* relation.

EXAMPLE :: Find NAMES OF ALL INSTRUCTORS IN "Comp Sci" Department ?

## A Sample Relational Model

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222     | Einstein    | Physics          | 95000         |
| 12121     | Wu          | Finance          | 90000         |
| 32343     | El Said     | History          | 60000         |
| 45565     | Katz        | Comp. Sci.       | 75000         |
| 98345     | Kim         | Elec. Eng.       | 80000         |
| 76766     | Crick       | Biology          | 72000         |
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 58583     | Califieri   | History          | 62000         |
| 83821     | Brandt      | Comp. Sci.       | 92000         |
| 15151     | Mozart      | Music            | 40000         |
| 33456     | Gold        | Physics          | 87000         |
| 76543     | Singh       | Finance          | 80000         |

(a) The *instructor* table

| <i>dept_name</i> | <i>building</i> | <i>budget</i> |
|------------------|-----------------|---------------|
| Comp. Sci.       | Taylor          | 100000        |
| Biology          | Watson          | 90000         |
| Elec. Eng.       | Taylor          | 85000         |
| Music            | Packard         | 80000         |
| Finance          | Painter         | 120000        |
| History          | Painter         | 50000         |
| Physics          | Watson          | 70000         |

(b) The *department* table

1) FROM

# Utilizing Relational DBMS: Lifecycle

- 1) 1. Design relational schema *(and tell the system about it)*
- 2) 2. Populate tables/relations
- 3) 3. Write queries to get information back from tables

# “Database Languages” Supported by DBMS Interface

- The DBMS interface supports the following types of languages:
- Data Definition Language (DDL)
  - Used to define the Database Schema
- Data Manipulation Language (DML)
  - Used to Retrieve / manipulate data



# “Database Languages” Supported by DBMS Interface

- The DBMS interface supports the following types of languages:
- Data Definition Language (DDL)
  - Used to define the Database Schema
- Data Manipulation Language (DML)
  - Used to Retrieve / manipulate data
- In practice, the DDL and DML are not two separate languages: instead they simply form parts of a single database language.
- For relational databases, the most popular database language is Structured Query Language (SQL)

# Structured Query Language (SQL)

- The standard to access/retrieve/manipulate data in a **relational database**
- Examples of a Data Definition Language (DDL) Component

```
create table department  
  (dept_name  char (20),  
   building   char (15),  
   budget     numeric (12,2));
```


- Examples of a Data Manipulation Language (DML) Component

```
select instructor.name  
from instructor  
where instructor.dept_name = 'History';
```

```
select instructor.ID, department.dept_name  
from instructor, department  
where instructor.dept_name= department.dept_name and  
      department.budget > 95000;
```

# SQL: The History

- IBM developed the original version of SQL, originally called Sequel, in the early 1970s.
- SQL clearly established itself as the standard relational database language.
- In 1986, ANSI and ISO published an SQL standard, called SQL-86.
- Follow-up standards
  - SQL-89
  - SQL-92
  - SQL:1999
  - SQL:2003
  - SQL:2006
  - SQL:2008
  - SQL:2011
  - SQL:2016



# SQL: Overview

- The SQL has two main parts:

**Data-definition language (DDL).** The SQL DDL provides commands for defining relation schemas, deleting relations, and modifying relation schemas.

**Data-manipulation language (DML).** The SQL DML provides the ability to query information from the database and to insert tuples into, delete tuples from, and modify tuples in the database.

# SQL: Data Definition

- Create Tables
- Domain Types
- Identify constraints (primary key / foreign key)

# Creating Tables

The general form of the create table command is:

```
create table r
  (A1 D1,
   A2 D2,
   ...,
   An Dn,
   ⟨integrity-constraint1⟩,
   ...,
   ⟨integrity-constraintk⟩);
```

where *r* is the name of the relation, each *A<sub>i</sub>* is the name of an attribute in the schema of relation *r*, and *D<sub>i</sub>* is the domain of attribute *A<sub>i</sub>*; that is, *D<sub>i</sub>* specifies the type of attribute *A<sub>i</sub>* along with optional constraints that restrict the set of allowed values for *A<sub>i</sub>*.

```
create table department
  (dept_name  varchar (20),
   building   varchar (15),
   budget     numeric (12,2),
   primary key (dept_name));
```

```
create table instructor
  (ID          varchar (5),
   name        varchar (20) not null,
   dept_name   varchar (20),
   salary      numeric (8,2),
   primary key (ID),
   foreign key (dept_name) references department);
```

# Domain Types

The SQL standard supports a variety of built-in types, including:

- **char(*n*)**: A fixed-length character string with user-specified length *n*. The full form, **character**, can be used instead.
- **varchar(*n*)**: A variable-length character string with user-specified maximum length *n*. The full form, **character varying**, is equivalent.
- **int**: An integer (a finite subset of the integers that is machine dependent). The full form, **integer**, is equivalent.
- **smallint**: A small integer (a machine-dependent subset of the integer type).
- **numeric(*p*, *d*)**: A fixed-point number with user-specified precision. The number consists of *p* digits (plus a sign), and *d* of the *p* digits are to the right of the decimal point. Thus, **numeric(3,1)** allows 44.5 to be stored exactly, but neither 444.5 nor 0.32 can be stored exactly in a field of this type.
- **real**, **double precision**: Floating-point and double-precision floating-point numbers with machine-dependent precision.
- **float(*n*)**: A floating-point number with precision of at least *n* digits.

# Integrity Constraints

- Types of integrity constraints
  - **primary key**  $(A_1, \dots, A_n)$
  - **foreign key**  $(A_m, \dots, A_n)$  **references**  $r$
  - **not null**
- SQL prevents any update to the database that violates an integrity constraint.



## SQL Queries: University Database (1/4)

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 12121     | Wu          | Finance          | 90000         |
| 15151     | Mozart      | Music            | 40000         |
| 22222     | Einstein    | Physics          | 95000         |
| 32343     | El Said     | History          | 60000         |
| 33456     | Gold        | Physics          | 87000         |
| 45565     | Katz        | Comp. Sci.       | 75000         |
| 58583     | Califieri   | History          | 62000         |
| 76543     | Singh       | Finance          | 80000         |
| 76766     | Crick       | Biology          | 72000         |
| 83821     | Brandt      | Comp. Sci.       | 92000         |
| 98345     | Kim         | Elec. Eng.       | 80000         |

**Figure 2.1** The *instructor* relation.

| <i>ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> |
|-----------|------------------|---------------|-----------------|-------------|
| 10101     | CS-101           | 1             | Fall            | 2017        |
| 10101     | CS-315           | 1             | Spring          | 2018        |
| 10101     | CS-347           | 1             | Fall            | 2017        |
| 12121     | FIN-201          | 1             | Spring          | 2018        |
| 15151     | MU-199           | 1             | Spring          | 2018        |
| 22222     | PHY-101          | 1             | Fall            | 2017        |
| 32343     | HIS-351          | 1             | Spring          | 2018        |
| 45565     | CS-101           | 1             | Spring          | 2018        |
| 45565     | CS-319           | 1             | Spring          | 2018        |
| 76766     | BIO-101          | 1             | Summer          | 2017        |
| 76766     | BIO-301          | 1             | Summer          | 2018        |
| 83821     | CS-190           | 1             | Spring          | 2017        |
| 83821     | CS-190           | 2             | Spring          | 2017        |
| 83821     | CS-319           | 2             | Spring          | 2018        |
| 98345     | EE-181           | 1             | Spring          | 2017        |

**Figure 2.7** The *teaches* relation.

## SQL Queries: University Database (2/4)

| <i>dept_name</i> | <i>building</i> | <i>budget</i> |
|------------------|-----------------|---------------|
| Biology          | Watson          | 90000         |
| Comp. Sci.       | Taylor          | 100000        |
| Elec. Eng.       | Taylor          | 85000         |
| Finance          | Painter         | 120000        |
| History          | Painter         | 50000         |
| Music            | Packard         | 80000         |
| Physics          | Watson          | 70000         |

**Figure 2.5** The *department* relation.

| <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> | <i>building</i> | <i>room_number</i> | <i>time_slot_id</i> |
|------------------|---------------|-----------------|-------------|-----------------|--------------------|---------------------|
| BIO-101          | 1             | Summer          | 2017        | Painter         | 514                | B                   |
| BIO-301          | 1             | Summer          | 2018        | Painter         | 514                | A                   |
| CS-101           | 1             | Fall            | 2017        | Packard         | 101                | H                   |
| CS-101           | 1             | Spring          | 2018        | Packard         | 101                | F                   |
| CS-190           | 1             | Spring          | 2017        | Taylor          | 3128               | E                   |
| CS-190           | 2             | Spring          | 2017        | Taylor          | 3128               | A                   |
| CS-315           | 1             | Spring          | 2018        | Watson          | 120                | D                   |
| CS-319           | 1             | Spring          | 2018        | Watson          | 100                | B                   |
| CS-319           | 2             | Spring          | 2018        | Taylor          | 3128               | C                   |
| CS-347           | 1             | Fall            | 2017        | Taylor          | 3128               | A                   |
| EE-181           | 1             | Spring          | 2017        | Taylor          | 3128               | C                   |
| FIN-201          | 1             | Spring          | 2018        | Packard         | 101                | B                   |
| HIS-351          | 1             | Spring          | 2018        | Painter         | 514                | C                   |
| MU-199           | 1             | Spring          | 2018        | Packard         | 101                | D                   |
| PHY-101          | 1             | Fall            | 2017        | Watson          | 100                | A                   |

**Figure 2.6** The *section* relation.

## SQL Queries: University Database (3/4)

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301          | BIO-101          |
| BIO-399          | BIO-101          |
| CS-190           | CS-101           |
| CS-315           | CS-101           |
| CS-319           | CS-101           |
| CS-347           | CS-101           |
| EE-181           | PHY-101          |

**Figure 2.3** The *prereq* relation.

| <i>course_id</i> | <i>title</i>               | <i>dept_name</i> | <i>credits</i> |
|------------------|----------------------------|------------------|----------------|
| BIO-101          | Intro. to Biology          | Biology          | 4              |
| BIO-301          | Genetics                   | Biology          | 4              |
| BIO-399          | Computational Biology      | Biology          | 3              |
| CS-101           | Intro. to Computer Science | Comp. Sci.       | 4              |
| CS-190           | Game Design                | Comp. Sci.       | 4              |
| CS-315           | Robotics                   | Comp. Sci.       | 3              |
| CS-319           | Image Processing           | Comp. Sci.       | 3              |
| CS-347           | Database System Concepts   | Comp. Sci.       | 3              |
| EE-181           | Intro. to Digital Systems  | Elec. Eng.       | 3              |
| FIN-201          | Investment Banking         | Finance          | 3              |
| HIS-351          | World History              | History          | 3              |
| MU-199           | Music Video Production     | Music            | 3              |
| PHY-101          | Physical Principles        | Physics          | 4              |

**Figure 2.2** The *course* relation.

## SQL Queries: University Database (4/4)

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>tot_cred</i> |
|-----------|-------------|------------------|-----------------|
| 00128     | Zhang       | Comp. Sci.       | 102             |
| 12345     | Shankar     | Comp. Sci.       | 32              |
| 19991     | Brandt      | History          | 80              |
| 23121     | Chavez      | Finance          | 110             |
| 44553     | Peltier     | Physics          | 56              |
| 45678     | Levy        | Physics          | 46              |
| 54321     | Williams    | Comp. Sci.       | 54              |
| 55739     | Sanchez     | Music            | 38              |
| 70557     | Snow        | Physics          | 0               |
| 76543     | Brown       | Comp. Sci.       | 58              |
| 76653     | Aoi         | Elec. Eng.       | 60              |
| 98765     | Bourikas    | Elec. Eng.       | 98              |
| 98988     | Tanaka      | Biology          | 120             |

**Figure 4.1** The *student* relation.

| <i>ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> | <i>grade</i> |
|-----------|------------------|---------------|-----------------|-------------|--------------|
| 00128     | CS-101           | 1             | Fall            | 2017        | A            |
| 00128     | CS-347           | 1             | Fall            | 2017        | A-           |
| 12345     | CS-101           | 1             | Fall            | 2017        | C            |
| 12345     | CS-190           | 2             | Spring          | 2017        | A            |
| 12345     | CS-315           | 1             | Spring          | 2018        | A            |
| 12345     | CS-347           | 1             | Fall            | 2017        | A            |
| 19991     | HIS-351          | 1             | Spring          | 2018        | B            |
| 23121     | FIN-201          | 1             | Spring          | 2018        | C+           |
| 44553     | PHY-101          | 1             | Fall            | 2017        | B-           |
| 45678     | CS-101           | 1             | Fall            | 2017        | F            |
| 45678     | CS-101           | 1             | Spring          | 2018        | B+           |
| 45678     | CS-319           | 1             | Spring          | 2018        | B            |
| 54321     | CS-101           | 1             | Fall            | 2017        | A-           |
| 54321     | CS-190           | 2             | Spring          | 2017        | B+           |
| 55739     | MU-199           | 1             | Spring          | 2018        | A-           |
| 76543     | CS-101           | 1             | Fall            | 2017        | A            |
| 76543     | CS-319           | 2             | Spring          | 2018        | A            |
| 76653     | EE-181           | 1             | Spring          | 2017        | C            |
| 98765     | CS-101           | 1             | Fall            | 2017        | C-           |
| 98765     | CS-315           | 1             | Spring          | 2018        | B            |
| 98988     | BIO-101          | 1             | Summer          | 2017        | A            |
| 98988     | BIO-301          | 1             | Summer          | 2018        | <i>null</i>  |

**Figure 4.2** The *takes* relation.

# SQL: Data Manipulation

- Basic query structure
- Major clauses in an SQL query
  - SELECT
  - FROM
  - WHERE
  - GROUP BY
  - HAVING

EXAMPLE :: Find NAMES OF ALL INSTRUCTORS IN "Comp Sci" Department ?

## SQL Queries: Basic Query Structure

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 12121 | Wu         | Finance    | 90000  |
| 15151 | Mozart     | Music      | 40000  |
| 22222 | Einstein   | Physics    | 95000  |
| 32343 | El Said    | History    | 60000  |
| 33456 | Gold       | Physics    | 87000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 58583 | Califieri  | History    | 62000  |
| 76543 | Singh      | Finance    | 80000  |
| 76766 | Crick      | Biology    | 72000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |

Figure 2.1 The *instructor* relation.

| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2017 |
| 10101 | CS-315    | 1      | Spring   | 2018 |
| 10101 | CS-347    | 1      | Fall     | 2017 |
| 12121 | FIN-201   | 1      | Spring   | 2018 |
| 15151 | MU-199    | 1      | Spring   | 2018 |
| 22222 | PHY-101   | 1      | Fall     | 2017 |
| 32343 | HIS-351   | 1      | Spring   | 2018 |
| 45565 | CS-101    | 1      | Spring   | 2018 |
| 45565 | CS-319    | 1      | Spring   | 2018 |
| 76766 | BIO-101   | 1      | Summer   | 2017 |
| 76766 | BIO-301   | 1      | Summer   | 2018 |
| 83821 | CS-190    | 1      | Spring   | 2017 |
| 83821 | CS-190    | 2      | Spring   | 2017 |
| 83821 | CS-319    | 2      | Spring   | 2018 |
| 98345 | EE-181    | 1      | Spring   | 2017 |

Figure 2.7 The *teaches* relation.

3) SELECT name

1) FROM instructor

2) WHERE dept\_name = 'Comp. Sci'

('Project' operation in relational algebra)

('Selection' operation in relational algebra)

# Basic Query Structure

- A typical SQL query has the form:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

- $A_i$  represents an attribute
  - $R_i$  represents a relation
  - $P$  is a predicate.
- The result of an SQL query is a relation.

## "Order of Execution" in SQL Queries

In general, the meaning of an SQL query can be understood as follows:

1. Generate a Cartesian product of the relations listed in the from clause.
2. Apply the predicates specified in the where clause on the result of Step 1.
3. For each tuple in the result of Step 2, output the attributes (or results of expressions) specified in the select clause.



# SELECT Clause

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra

- Example: find the names of all instructors:

**select** *name*  
**from** *instructor*

- NOTE: SQL names are case insensitive
  - E.g., *Name*  $\equiv$  *NAME*  $\equiv$  *name*
  - Some people use upper case wherever bold font is used.

|    | name       |
|----|------------|
| 1  | Srinivasan |
| 2  | Wu         |
| 3  | Mozart     |
| 4  | Einstein   |
| 5  | El Said    |
| 6  | Gold       |
| 7  | Katz       |
| 8  | Califieri  |
| 9  | Singh      |
| 10 | Crick      |
| 11 | Brandt     |
| 12 | Kim        |

## SELECT Clause: distinct and all Keywords

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after select.
- Find the department names of all instructors, and remove duplicates

```
select distinct dept_name  
from instructor
```

|   | dept_name  |
|---|------------|
| 1 | Biology    |
| 2 | Comp. Sci. |
| 3 | Elec. Eng. |
| 4 | Finance    |
| 5 | History    |
| 6 | Music      |
| 7 | Physics    |

| dept_name         |
|-------------------|
| <u>Comp. Sci.</u> |
| Finance           |
| Music             |
| Physics           |
| History           |
| Physics           |
| <u>Comp. Sci.</u> |
| History           |
| Finance           |
| Biology           |
| Comp. Sci.        |
| Elec. Eng.        |

- The keyword **all** specifies that duplicates should not be removed.

```
select all dept_name  
from instructor
```

# SELECT Clause

- An asterisk in the select clause denotes “all attributes”
- An attribute can also be renamed using ***as***

**select** \*  
**from** *instructor*

|    | ID    | name       | dept_name  | salary   |
|----|-------|------------|------------|----------|
| 1  | 10101 | Srinivasan | Comp. Sci. | 65000.00 |
| 2  | 12121 | Wu         | Finance    | 90000.00 |
| 3  | 15151 | Mozart     | Music      | 40000.00 |
| 4  | 22222 | Einstein   | Physics    | 95000.00 |
| 5  | 32343 | El Said    | History    | 60000.00 |
| 6  | 33456 | Gold       | Physics    | 87000.00 |
| 7  | 45565 | Katz       | Comp. Sci. | 75000.00 |
| 8  | 58583 | Califieri  | History    | 62000.00 |
| 9  | 76543 | Singh      | Finance    | 80000.00 |
| 10 | 76766 | Crick      | Biology    | 72000.00 |
| 11 | 83821 | Brandt     | Comp. Sci. | 92000.00 |
| 12 | 98345 | Kim        | Elec. Eng. | 80000.00 |

**select** name as *nom*  
**from** *instructor*

|    | nom        |
|----|------------|
| 1  | Srinivasan |
| 2  | Wu         |
| 3  | Mozart     |
| 4  | Einstein   |
| 5  | El Said    |
| 6  | Gold       |
| 7  | Katz       |
| 8  | Califieri  |
| 9  | Singh      |
| 10 | Crick      |
| 11 | Brandt     |
| 12 | Kim        |

## SELECT Clause: Operators

- The **select** clause can contain arithmetic expressions involving the operation, +, −, \*, and /, and operating on constants or attributes of tuples.

**select** *ID, name, salary/12 as monthly\_salary*  
**from** instructor

|    | ID    | name       | monthly_salary |
|----|-------|------------|----------------|
| 1  | 10101 | Srinivasan | 5416.666666    |
| 2  | 12121 | Wu         | 7500.000000    |
| 3  | 15151 | Mozart     | 3333.333333    |
| 4  | 22222 | Einstein   | 7916.666666    |
| 5  | 32343 | El Said    | 5000.000000    |
| 6  | 33456 | Gold       | 7250.000000    |
| 7  | 45565 | Katz       | 6250.000000    |
| 8  | 58583 | Califieri  | 5166.666666    |
| 9  | 76543 | Singh      | 6666.666666    |
| 10 | 76766 | Crick      | 6000.000000    |
| 11 | 83821 | Brandt     | 7666.666666    |
| 12 | 98345 | Kim        | 6666.666666    |

## 2) WHERE Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```

|   | name       |
|---|------------|
| 1 | Srinivasan |
| 2 | Katz       |
| 3 | Brandt     |

## WHERE Clause: Logical Connectives

- SQL allows the use of the logical connectives **and**, **or**, and **not**
- The operands of the logical connectives can be expressions involving the comparison operators **<**, **<=**, **>**, **>=**, **=**, and **<>**. (**!=**)
- Comparisons can be applied to results of arithmetic expressions
- To find all instructors in Comp. Sci. dept with salary > 70000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and salary > 70000
```

|   | name   |
|---|--------|
| 1 | Katz   |
| 2 | Brandt |

### 3) FROM Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *instructor X teaches*
  - ↪ select \***
    - 1, from instructor, teaches**
      - generates every possible instructor – teaches pair, with all attributes from both relations.

# FROM Clause: Example

**select \***  
**from** *instructor, teaches*

- Partial Result of the Cartesian Product:

|    | ID    | name       | dept_name  | salary   | ID    | course_id | sec_id | semester | year |
|----|-------|------------|------------|----------|-------|-----------|--------|----------|------|
| 1  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 10101 | CS-101    | 1      | Fall     | 2017 |
| 2  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 10101 | CS-315    | 1      | Spring   | 2018 |
| 3  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 10101 | CS-347    | 1      | Fall     | 2017 |
| 4  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 12121 | FIN-201   | 1      | Spring   | 2018 |
| 5  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 15151 | MU-199    | 1      | Spring   | 2018 |
| 6  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 22222 | PHY-101   | 1      | Fall     | 2017 |
| 7  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 32343 | HIS-351   | 1      | Spring   | 2018 |
| 8  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 45565 | CS-101    | 1      | Spring   | 2018 |
| 9  | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 45565 | CS-319    | 1      | Spring   | 2018 |
| 10 | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 76766 | BIO-101   | 1      | Summer   | 2017 |
| 11 | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 76766 | BIO-301   | 1      | Summer   | 2018 |
| 12 | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 83821 | CS-190    | 1      | Spring   | 2017 |
| 13 | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 83821 | CS-190    | 2      | Spring   | 2017 |
| 14 | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 83821 | CS-319    | 2      | Spring   | 2018 |
| 15 | 10101 | Srinivasan | Comp. Sci. | 65000.00 | 98345 | EE-181    | 1      | Spring   | 2017 |
| 16 | 12121 | Wu         | Finance    | 90000.00 | 10101 | CS-101    | 1      | Fall     | 2017 |
| 17 | 12121 | Wu         | Finance    | 90000.00 | 10101 | CS-315    | 1      | Spring   | 2018 |
| 18 | 12121 | Wu         | Finance    | 90000.00 | 10101 | CS-347    | 1      | Fall     | 2017 |
| 19 | 12121 | Wu         | Finance    | 90000.00 | 12121 | FIN-201   | 1      | Spring   | 2018 |



## FROM Clause: Example

- Cartesian product is not very useful directly, but it is useful when combined with where-clause condition (selection operation in relational algebra).
- Find the names of all instructors who have taught some course and the course\_id

3) **select** *name, course\_id*  
1) **from** *instructor, teaches*  
2) **where** *instructor.ID* = *teaches.ID*

|    | name       | course_id |
|----|------------|-----------|
| 1  | Srinivasan | CS-101    |
| 2  | Srinivasan | CS-315    |
| 3  | Srinivasan | CS-347    |
| 4  | Wu         | FIN-201   |
| 5  | Mozart     | MU-199    |
| 6  | Einstein   | PHY-101   |
| 7  | El Said    | HIS-351   |
| 8  | Katz       | CS-101    |
| 9  | Katz       | CS-319    |
| 10 | Crick      | BIO-101   |
| 11 | Crick      | BIO-301   |
| 12 | Brandt     | CS-190    |
| 13 | Brandt     | CS-190    |
| 14 | Brandt     | CS-319    |
| 15 | Kim        | EE-181    |

# String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring.
  - underscore ( \_ ). The \_ character matches any character.
- Find the names of all instructors whose name includes the substring “ri”.

```
select name  
from instructor  
where name like '%ri%'
```

| <b>name</b> |
|-------------|
| Srinivasan  |
| Califieri   |
| Crick       |

# String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring.
  - underscore ( \_ ). The \_ character matches any character.
- Find the names of all instructors whose name includes the substring “ri”.

```
select name  
from instructor  
where name like '%ri%'
```

| name       |
|------------|
| Srinivasan |
| Califieri  |
| Crick      |

- Match the string “100%”

```
like '100 \%' escape '\'
```

in that above we use backslash (\) as the escape character.

# String Operations

- Patterns are case sensitive.
- Pattern matching examples:
  - 'Intro%' matches any string beginning with “Intro”.
  - '%Comp%' matches any string containing “Comp” as a substring.
  - '\_\_\_' matches any string of exactly three characters.
  - '\_\_\_%' matches any string of at least three characters.
- SQL supports a variety of string operations such as
  - concatenation (using “||”)
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.

| course_id |
|-----------|
| CS-101    |
| CS-102    |
| ...       |

| course_id |
|-----------|
| CS-101    |
| CS-312    |
| ...       |

## Set Operations

- Find courses that ran in Fall 2017 or in Spring 2018  
 (select course\_id from section where semester = 'Fall' and year = 2017)  
union  
 (select course\_id from section where semester = 'Spring' and year = 2018)
- Find courses that ran in Fall 2017 and in Spring 2018  
 (select course\_id from section where semester = 'Fall' and year = 2017)  
 ✓ intersect  
 (select course\_id from section where semester = 'Spring' and year = 2018)
- Find courses that ran in Fall 2017 but not in Spring 2018  
 (select course\_id from section where semester = 'Fall' and year = 2017)  
except  
 (select course\_id from section where semester = 'Spring' and year = 2018)

# Set Operations

- Set operations **union**, **intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the
  - **union all**,
  - **intersect all**
  - **except all**.

# Null Values

- It is possible for tuples to have a null value, denoted by **null**, for some of their attributes
- **null** signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving **null** is **null**
  - Example: 5 + null returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.  

```
select name  
from instructor  
where salary is null
```
- The predicate **is not null** succeeds if the value on which it is applied is not null.

# Null Values

- SQL treats as unknown the result of any comparison involving a null value (other than predicates **is null** and **is not null**).
  - Example: 5 < null or **null <> null** or **null = null**
- The predicate in a **where** clause can involve Boolean operations (**and**, **or**, **not**); thus the definitions of the Boolean operations need to be extended to deal with the value **unknown**.
  - **and** : (*true and unknown*) = *unknown*,  
          (*false and unknown*) = *false*,  
          (*unknown and unknown*) = *unknown*
  - **or** : (*unknown or true*) = *true*,  
          (*unknown or false*) = *unknown*  
          (*unknown or unknown*) = *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*



# Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values

# Aggregate Functions

- Find the average salary of instructors in the Computer Science department

# Aggregate Functions

- Find the average salary of instructors in the Computer Science department

3) SELECT avg(salary)

1) FROM instructor

2) WHERE dept\_name = 'Comp Sci'

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 12121 | Wu         | Finance    | 90000  |
| 15151 | Mozart     | Music      | 40000  |
| 22222 | Einstein   | Physics    | 95000  |
| 32343 | El Said    | History    | 60000  |
| 33456 | Gold       | Physics    | 87000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 58583 | Califieri  | History    | 62000  |
| 76543 | Singh      | Finance    | 80000  |
| 76766 | Crick      | Biology    | 72000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |

Figure 2.1 The *instructor* relation.

# Aggregate Functions

- Find the average salary of instructors in the Computer Science department

```
select avg (salary)  
from instructor  
where dept_name= 'Comp. Sci.';
```

| avg (salary)      |
|-------------------|
| 77333.33333333333 |

# Aggregate Functions

- Find the total number of instructors who teach a course in the Spring 2018 semester

3) SELECT <sup>distinct</sup> count ( ID )  
 1) FROM teaches  
 2) WHERE semester = 'Spring' and year = 2018

| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2017 |
| 10101 | CS-315    | 1      | Spring   | 2018 |
| 10101 | CS-347    | 1      | Fall     | 2017 |
| 12121 | FIN-201   | 1      | Spring   | 2018 |
| 15151 | MU-199    | 1      | Spring   | 2018 |
| 22222 | PHY-101   | 1      | Fall     | 2017 |
| 32343 | HIS-351   | 1      | Spring   | 2018 |
| 45565 | CS-101    | 1      | Spring   | 2018 |
| 45565 | CS-319    | 1      | Spring   | 2018 |
| 76766 | BIO-101   | 1      | Summer   | 2017 |
| 76766 | BIO-301   | 1      | Summer   | 2018 |
| 83821 | CS-190    | 1      | Spring   | 2017 |
| 83821 | CS-190    | 2      | Spring   | 2017 |
| 83821 | CS-319    | 2      | Spring   | 2018 |
| 98345 | EE-181    | 1      | Spring   | 2017 |

Figure 2.7 The *teaches* relation.

# Aggregate Functions

- Find the total number of instructors who teach a course in the Spring 2018 semester

```
select count (distinct ID)  
from teaches  
where semester = 'Spring' and year = 2018;
```

| count (distinct ID) |
|---------------------|
| 6                   |

# Aggregate Functions

- Find the number of tuples in the *course* relation

| <i>course_id</i> | <i>title</i>               | <i>dept_name</i> | <i>credits</i> |
|------------------|----------------------------|------------------|----------------|
| BIO-101          | Intro. to Biology          | Biology          | 4              |
| BIO-301          | Genetics                   | Biology          | 4              |
| BIO-399          | Computational Biology      | Biology          | 3              |
| CS-101           | Intro. to Computer Science | Comp. Sci.       | 4              |
| CS-190           | Game Design                | Comp. Sci.       | 4              |
| CS-315           | Robotics                   | Comp. Sci.       | 3              |
| CS-319           | Image Processing           | Comp. Sci.       | 3              |
| CS-347           | Database System Concepts   | Comp. Sci.       | 3              |
| EE-181           | Intro. to Digital Systems  | Elec. Eng.       | 3              |
| FIN-201          | Investment Banking         | Finance          | 3              |
| HIS-351          | World History              | History          | 3              |
| MU-199           | Music Video Production     | Music            | 3              |
| PHY-101          | Physical Principles        | Physics          | 4              |

3) SELECT count ( \* )

1) FROM *course*

2) WHERE

**Figure 2.2** The *course* relation.

# Aggregate Functions

- Find the number of tuples in the *course* relation

```
select count (*)  
from course;
```

| count (*) |
|-----------|
| 13        |