

Database Systems

(CS 355 / CE 373)

Dr. Umer Tariq
Assistant Professor,
Dhanani School of Science & Engineering,
Habib University

Acknowledgements

- Many slides have been borrowed from the official lecture slides accompanying the textbook:

Database System Concepts, (2019), Seventh Edition,
Avi Silberschatz, Henry F. Korth, S. Sudarshan
McGraw-Hill, ISBN 9780078022159

The original lecture slides are available at:

<https://www.db-book.com/>

- Some of the slides have been borrowed from the lectures by Dr. Immanuel Trummer (Cornell University). Available at: (www.itrummer.org)

Outline: Week 12

- Interaction of Transaction Isolation and Atomicity Properties
 - Recoverable Schedules
 - Cascading Rollback
 - Cascadeless Schedules
- Transaction Isolation Levels
- Concurrency-Control Mechanisms
 - Locks

DBMS: The Concept of Transactions

DB

Accounts

| Number | Balance | C-ID |
|--------|---------|------|
| | | |

Account(C-ID, Number, Balance)

Application Code.

TRANSFER

On Button Clicked OK

(SELECT Balance Read(A)

FROM Accounts

WHERE Number = '...1245'

⇒ BalanceVar

BalanceVar = BalanceVar - 50

(Update Accounts Write(A)

Set Balance = BalanceVar

WHERE Number = '...1245'

(SELECT Balance Read(B)

FROM Accounts

WHERE Number = '...5679'

(BalanceVar
Balance = BalanceVar + 50

(Update Accounts Write(B)

Set Balance = BalanceVar

WHERE Number = '...5679'

UI

| Name | Transfer |
|------------------------|----------|
| Current Account - 1245 | 10,000 |
| Savings Account - 5679 | 50,000 |

From

----- 1245

To

5679

Amount 50 OK

DBMS: The Concept of Transactions

TRANSFER (Transfer. 181)

BEGIN TRANSACTION

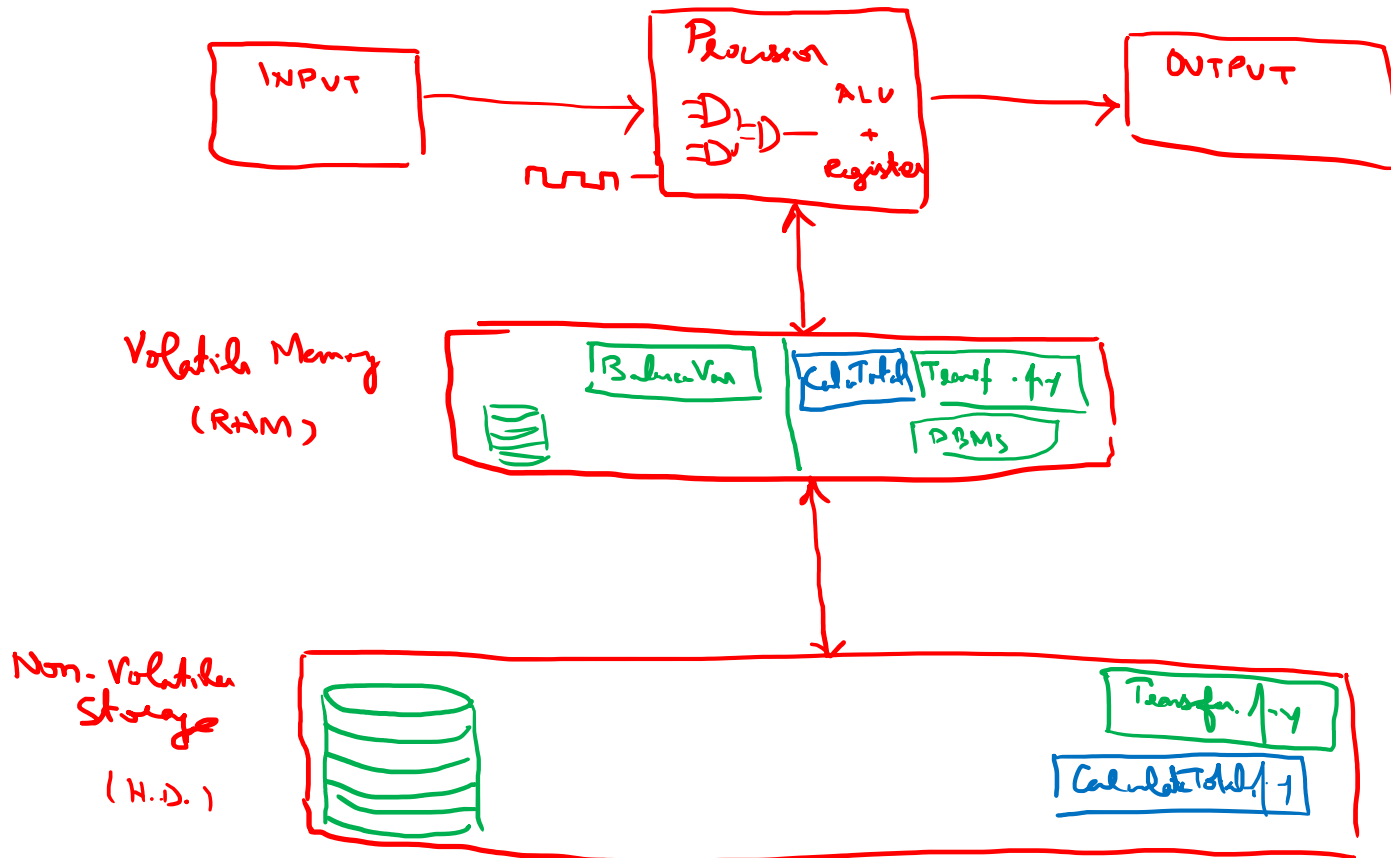
1. Read (A)
2. $A = A - 50$
3. Write (A)

4. Read (B)
5. $B = B + 50$
6. Write (B)

END TRANSACTION

DBMS: The Concept of Transactions

Draw a "COMPUTER"?



DBMS: The Concept of Transactions

- Transaction
 - A transaction is unit of program execution that consists of multiple database operations but appears as a single, indivisible unit from the point of view of the database user/application.
 - A transaction executes in its entirety or not at all.
- Example
 - A transaction to transfer Rs. 50 from account A to account B


```
Ti: read(A);  
      A := A - 50;  
      write(A);  
      read(B);  
      B := B + 50;  
      write(B).
```

Transaction Properties: ACID

- A transaction must have the following four properties:
 - Atomicity,
 - Consistency,
 - Isolation,
 - Durability.
- These form the acronym **ACID** properties.

Transaction States

- **Active**
 - the initial state; the transaction stays in this state while it is executing.
- **Partially committed**
 - after the final statement has been executed.
- **Failed**
 - after the discovery that normal execution can no longer proceed.
- **Aborted**
 - after the transaction has been rolled back and the database has been restored to its state prior to the start of the transaction.
- **Committed**
 - after successful completion.



TRANSFER:

1. **read(A)**
2. $A := A - 50$
3. **write(A)**
4. **read(B)**
5. $B := B + 50$
6. **write(B)**

Send-email()

Serial vs Concurrent Execution of Transactions

| T_1 | T_2 |
|---|---|
| <code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code> | <code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code> |

| T_1 |
|---|
| <code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code> |

| T_2 |
|---|
| <code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code> |

| T_1 | T_2 |
|---|---|
| <code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code> | <code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code> |

- Restricting ourselves to executing transactions serially (i.e. one after the other) makes it easy to achieve isolation among transactions.
- However, concurrent execution of transactions provides significant performance benefits:
 - Increased throughput
 - Reduced average response times

Concurrent Schedule: Consistent vs Inconsistent State

$A = 100, B = 100$

| T_1 | T_2 |
|---|--|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit |

$A = 45, B = 153$

| T_1 |
|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit |

$var A_1 = 100$
 $var A_2 = 50$
 $DB-A = 50$
 $var A = 50$
 $temp = 5$
 $var A_1 = 45$
 $DB-A = 45$
 $var B_1 = 100$
 $var B_2 = 150$
 $DB-B = 150$
 $var B_2 = 155$
 $DB-B = 155$

| T_2 |
|--|
| read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit |

$var A_1 = 100$
 $var A_2 = 50$
 $var A_2 = 100$
 $temp = 10$
 $var A_2 = 90$
 $DB-A = 90$
 $var B_2 = 100$
 $DB-A = 50$
 $var B_1 = 100$
 $var B_1 = 150$
 $DB-B = 150$
 $var B_2 = 110$
 $DB-B = 110$

$A = 100, B = 100$

| T_1 | T_2 |
|---|--|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit |

$A = 50, B = 110$

Serializable Schedule

- A serial schedule is always consistent, i.e. it maintains the consistent state of the database.
- However, the same cannot be guaranteed for a concurrent schedule.
- If a concurrent schedule can be shown to have the same effect as a serial schedule, (in other words it is shown to be equivalent to a serial schedule), then it can ensure the consistency of the database.
- Such a concurrent schedule is called a **serializable schedule**.

| T_1 |
|---|
| <code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code> |

| T_2 |
|---|
| <code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code> |

| T_1 | T_2 |
|---|---|
| <code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code> | <code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code> |

| T_1 | T_2 |
|---|---|
| <code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code> | <code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code> |

Serializable Schedules: Equivalence

| T_1 | T_2 |
|--|---|
| read(A) $A := A - 50$ write(A) | |
| | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) |
| read(B) $B := B + 50$ write(B) commit | |
| | read(B) $B := B + temp$ write(B) commit |

Figure 17.4 Schedule 3—a concurrent schedule equivalent to schedule 1.

**E
Q
U
I
V
A
L
E
N
T

T
O**

| T_1 | T_2 |
|--|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit | |
| | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit |

Figure 17.2 Schedule 1—a serial schedule in which T_1 is followed by T_2 .

If a concurrent schedule can be shown to be equivalent to a serial schedule, we conclude that this concurrent schedule maintains the consistency of the database.

Conflict-Equivalent Schedules

- If a schedule S can be transformed into a schedule S' by a series of swaps of nonconflicting instructions, we say that S and S' are **conflict equivalent**.

| T_1 | T_2 |
|-----------------------------|-----------------------------|
| read(A) write(A) | read(A) write(A) |
| read(B) write(B) | |
| | read(B) write(B) |

E
Q
U
I
V
A
L
E
N
T

T
O

| T_1 | T_2 |
|--|--|
| read(A) write(A) read(B) write(B) | read(A) write(A) read(B) write(B) |
| | |

Transaction Isolation and Atomicity

- So far, we have studied concurrent schedules while assuming implicitly that there are no transaction failures
- We now address the effect of transaction failures during concurrent execution of transactions.

Nonrecoverable Schedules



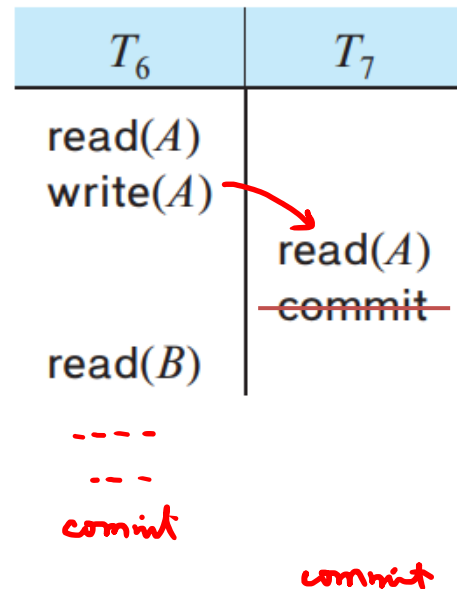
- Consider the partial schedule shown:
 - We call this a partial schedule because we have not included a commit or abort operation for T_6 .
- Notice the following:
 - T_7 has read the value of A written by T_6 . Therefore, we say that T_7 is dependent on T_6 .
 - T_7 commits immediately after executing the $read(A)$ instruction.
 - T_7 commits while T_6 is still in the active state.
- Suppose that T_6 fails after the $read(B)$ statement.
 - We must abort T_7 to ensure atomicity of T_6 (as any effect of value of A written by T_6 must be rolled back.)
 - However, T_7 has already been committed and cannot be aborted.
 - Thus, we have a situation where it is impossible to recover from the failure of T_6 .
 - This is an example of nonrecoverable schedule.

| T_6 | T_7 |
|-------------------------|--|
| $read(A)$ $write(A)$ | <div> $read(A)$ $commit$ </div> |
| $\times read(B)$ | |

...
...
...
commit

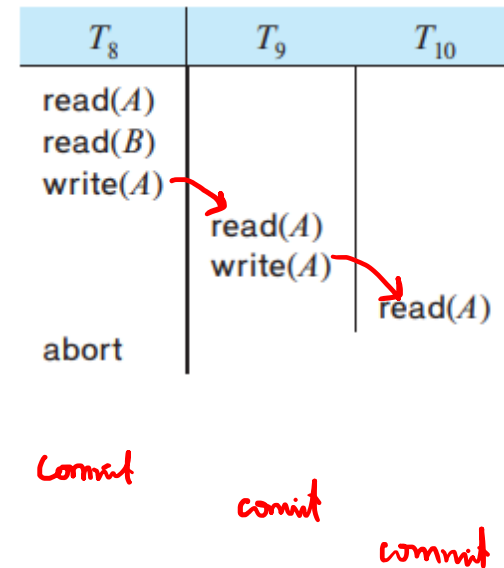
Recoverable Schedules

- A recoverable schedule is one where
 - For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the commit operation of T_j
- For the example shown here to be recoverable, T_7 would have to delay committing until after T_6 commits.



Cascading Rollback

- Even if a schedule is recoverable, to recover correctly from the failure of a transaction T_i , we may have to roll back several transactions.
- Consider the partial schedule shown here:
 - Transaction T_8 writes a value of A that is read by transaction T_9 . Transaction T_9 writes a value of A that is read by transaction T_{10} .
 - Suppose that, at this point (indicated as *abort*) T_8 fails.
 - T_8 must be rolled back. Since T_9 is dependent on T_8 , T_9 must be rolled back. Since T_{10} is dependent on T_9 , T_{10} must be rolled back.
 - This phenomenon, in which a single transaction failure leads to a series of transaction rollbacks, is called **cascading rollback**



Concurrent Schedules:

Serializable? Recoverable? Cascadeless?

Cascadeless Schedules

- Cascading rollback is undesirable, since it leads to the undoing of a significant amount of work.
- It is desirable to restrict the schedules to those where cascading rollbacks cannot occur.
-
- Such schedules are called **cascadeless schedules**.
- Formally,
 - a **cascadeless schedule** is one where, for each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- Note that a cascadeless schedule is also a recoverable schedule.

| T_8 | T_9 | T_{10} |
|--------------------------------|---------------------------------|--------------------|
| read(A) read(B) write(A) | read(A) write(A) | read(A) |

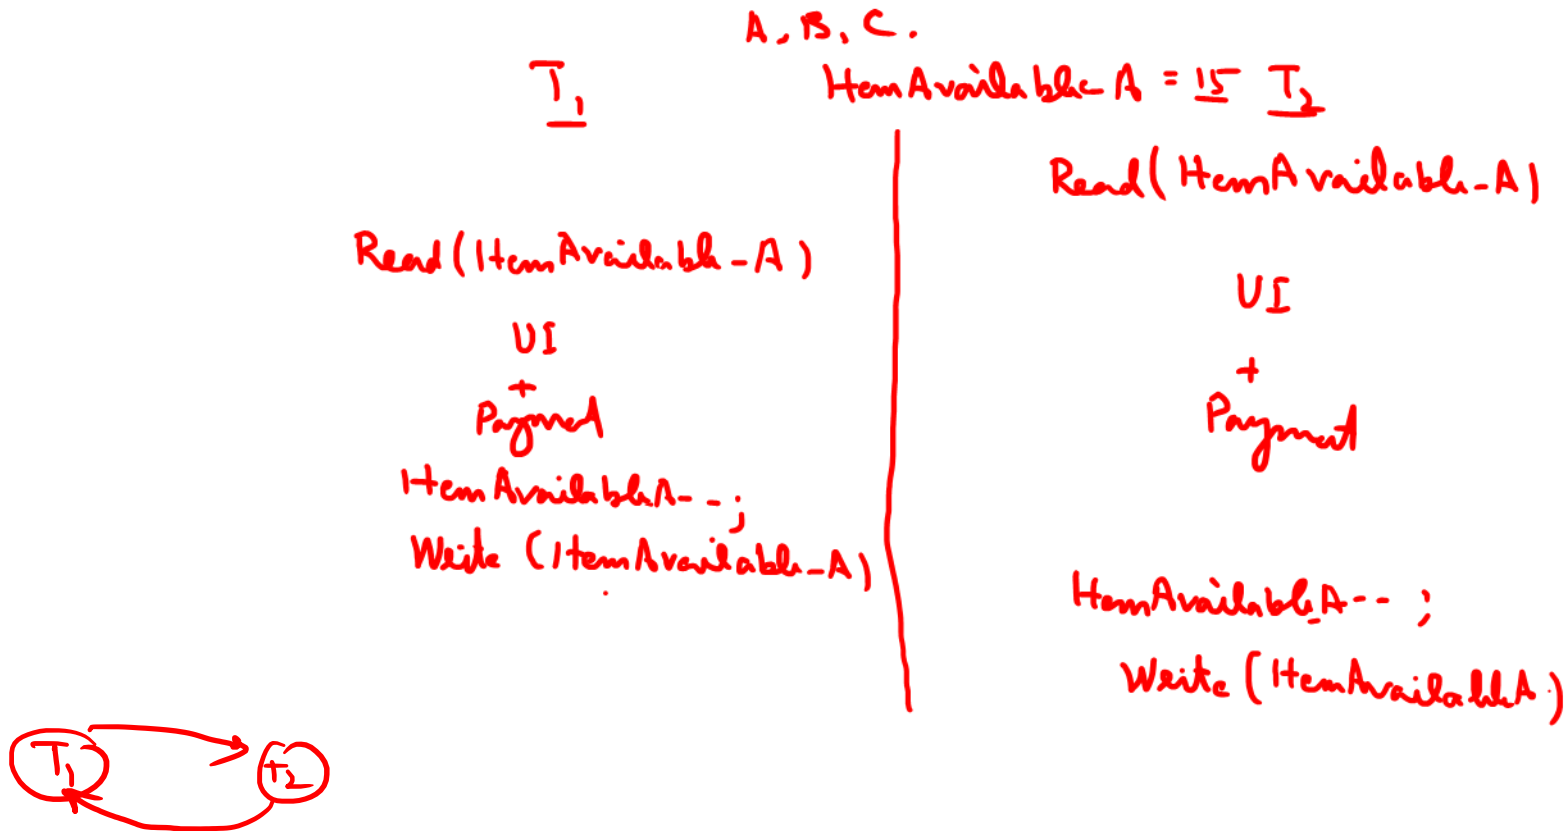
commit

read(A)
write(A)

commit
read(A)

Transaction Isolation Levels: Real-World Applicability

- Online Store Example (with two customers trying to buy the same item).



Transaction Isolation Levels: Background

- Transaction isolation (and serializability) is a useful concept because it allows programmers to ignore issues related to concurrency when they code transactions
- However, the protocols used to ensure serializability of concurrent transaction executions may allow too little concurrency for certain applications.
- In these cases, weaker levels of isolation/consistency are used.
 - The use of weaker levels of isolation places additional burdens on programmers for ensuring database correctness.
 - An application designer may decide to accept a weaker isolation level in order to improve system performance.

Transaction Isolation Levels: SQL Standard

The SQL standard allows a transaction to specify its isolation level:

- ✓ • **Serializable**
 - usually ensures serializable execution.
- **Repeatable read**
 - allows only committed data to be read and further requires that, between two reads of a data item by a transaction, no other transaction is allowed to update it.
- **Read committed**
 - allows only committed data to be read, but does not require repeatable reads.
- **Read uncommitted**
 - allows uncommitted data to be read. It is the lowest isolation level allowed by SQL.
- All the isolation levels above additionally disallow dirty writes: They disallow writes to a data item that has already been written by another transaction that has not yet committed or aborted.

Transaction Isolation Levels: SQL Standard

- Many database systems run, by default, at the read committed isolation level.
- In SQL, it is possible to set the isolation level explicitly. For example, through the statement:

```
set transaction isolation level serializable
```

Review: Concurrent Execution of Transactions: Role of Concurrency-Control Schemes

- Concurrency-control schemes
 - Mechanisms to achieve isolation among concurrently-executing transactions
 - Mechanisms to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database
- Will study these schemes after studying the notion of 'correctness of concurrent executions'

| T_1 | T_2 |
|--|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit |

| T_1 | T_2 |
|--|---|
| read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit |

Transaction Isolation and Concurrency Control Schemes

- Concurrency-control protocols of DBMS allow the concurrent execution of transaction in such a manner that the resulting transaction schedules are
 - Serializable
 - Recoverable
 - Cascadeless
- Concurrency control protocols (generally) do not examine the precedence graph as it is being created
 - Instead a protocol imposes a discipline that avoids non-serializable schedules.
- Different concurrency control protocols provide different tradeoffs between the amount of concurrency they allow and the amount of overhead that they incur.
- Tests for serializability help us understand why a concurrency-control protocol of DBMS is correct.

Concurrency Control Schemes: Overview

- There are various concurrency-control schemes that a DBMS can use to ensure that even when multiple transactions are executed concurrently, only acceptable transaction schedules are generated.
- Some of the main concurrency control schemes can be categorized based on their use of the following concepts
 - Locks
 - Timestamps
 - Multiple Versions of Data and Snapshot Isolation

A

B



Concurrency-Control Schemes: Locks

- One way to ensure isolation is to require that data items be accessed in a mutually exclusive manner
 - While one transaction is accessing a data item, no other transaction can modify that data item.
- The most common method used to implement this requirement is to allow a transaction to access a data item only if it is holding a **lock** on that item.
 - A lock is a mechanism which provides access privileges to a transaction on a shared resource,
 - While the lock is acquired, no other transaction can access that resource.
 - In order to release the resource, i.e. to make it available to other transactions, a transaction has to **unlock** the resource.

✓
 T_1 : lock-X(B);
✓ read(B);
✓ $B := B - 50$;
✓ write(B);
unlock(B);
lock-X(A);
✓ read(A);
✓ $A := A + 50$;
✓ write(A);
unlock(A).

✓
 T_2 : lock-S(A);
read(A);
unlock(A);
lock-S(B);
read(B);
unlock(B);
display(A + B).

Concurrency Control Schemes: Locks

| T_1 | T_2 | concurrency-control manager |
|---|--|-----------------------------|
| lock-X(B) | | grant-X(B, T_1) |
| read(B) $B := B - 50$ write(B) unlock(B) | lock(B) | x x |
| | lock-S(A) | grant-S(A, T_2) |
| | read(A) unlock(A) lock-S(B) | |
| | read(B) unlock(B) display($A + B$) | grant-S(B, T_2) |
| lock-X(A) | | |
| read(A) $A := A + 50$ write(A) unlock(A) | | grant-X(A, T_1) |

T_3

lock(B)

wait

lock(B)

Concurrency Control Schemes: Lock Modes

- There could be various possible modes in which the data item is locked. For example:
 - **Shared**. If a transaction T_i has obtained a **shared-mode** lock (**S**) on item Q , then T_i can read, but cannot write Q .
 - **Exclusive**. If a transaction T_i has obtained an **exclusive-mode** lock (**X**) on item Q , then T_i can both read and write Q .
- Every transaction must first request the **concurrency-control manager** for a lock in an appropriate mode on data item Q depending on the type of operations that it will perform on Q .
- The transaction can proceed with the operation only after the concurrency-control manager grants the lock to the transaction.
- The use of these two lock modes allow multiple transactions to read a data item but limits write access to just one instruction at a time.

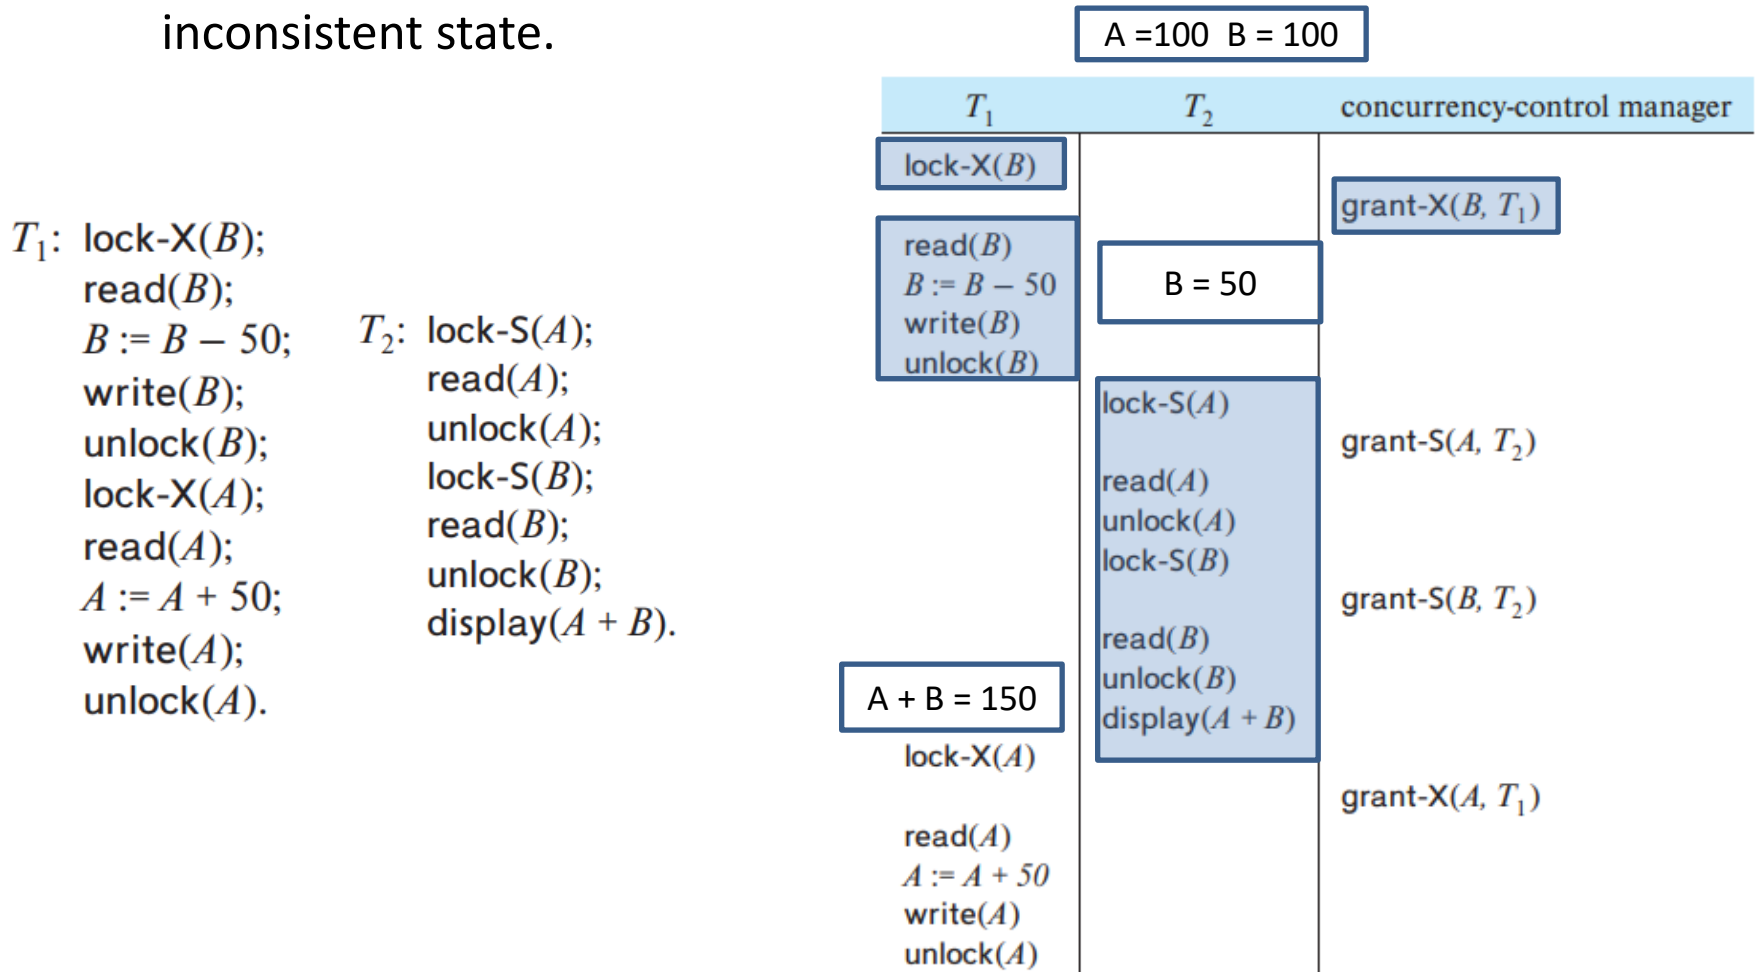
Concurrency Control Schemes: Compatibility of Lock Modes

- Given a set of lock modes, we can define a **compatibility function** on them as follows:
 - Let A and B represent arbitrary lock modes.
 - Suppose that a transaction T_i requests a lock of mode A on item Q on which transaction T_j ($T_i \neq T_j$) currently holds a lock of mode B .
 - If transaction T_i can be granted a lock on Q immediately, in spite of the presence of the mode B lock, then we say mode A is **compatible** with mode B .
- Such a compatibility function between a set of lock modes can be represented conveniently by a matrix (Lock-compatibility matrix)
 - An element **comp**(A, B) of the matrix has the value true if and only if mode A is compatible with mode B .
 - Compatibility matrix of shared-mode (S) and exclusive-mode (X) locks:

| | S | X |
|---|-------|-------|
| S | true | false |
| X | false | false |

Analysis of Transaction Schedules with Locks

- Despite the acquisition of locks the following schedule results in an inconsistent state.



Analysis of Transaction Schedules with Locks

- With unlocking delayed, transactions T_3 and T_4 lead to a consistent schedule
 - A Locking Protocol dictates when a transaction may lock and unlock each of the data item. The choice of a locking protocol restricts the number of possible transaction schedules.

T_1 : lock-X(B);
read(B);
 $B := B - 50$;
write(B);
unlock(B);
lock-X(A);
read(A);
 $A := A + 50$;
write(A);
unlock(A).

T_2 : lock-S(A);
read(A);
unlock(A);
lock-S(B);
read(B);
unlock(B);
display($A + B$).

T_3 : lock-X(B);
read(B);
 $B := B - 50$;
write(B);
lock-X(A);
read(A);
 $A := A + 50$;
write(A);
unlock(B);
unlock(A).

T_4 : lock-S(A);
read(A);
lock-S(B);
read(B);
display($A + B$);
unlock(A);
unlock(B).

The Two-Phase Locking Protocol

- One protocol that ensures serializability is the two-phase locking protocol.
- This protocol requires that each transaction issue lock and unlock requests in two phases:
 - Growing phase. A transaction may obtain locks, but may not release any lock.
 - Shrinking phase. A transaction may release locks, but may not obtain any new locks.
- Initially, a transaction is in the growing phase. The transaction acquires locks as needed. Once the transaction releases a lock, it enters the shrinking phase, and it can issue no more lock requests

```
T1: lock-X(B);  
    read(B);  
    B := B - 50;  
    write(B);  
    unlock(B);  
    lock-X(A);  
    read(A);  
    A := A + 50;  
    write(A);  
    unlock(A).
```

```
T2: lock-S(A);  
    read(A);  
    unlock(A);  
    lock-S(B);  
    read(B);  
    unlock(B);  
    display(A + B).
```

```
T3: lock-X(B);  
    read(B);  
    B := B - 50;  
    write(B);  
    lock-X(A);  
    read(A);  
    A := A + 50;  
    write(A);  
    unlock(B);  
    unlock(A).
```

```
T4: lock-S(A);  
    read(A);  
    lock-S(B);  
    read(B);  
    display(A + B);  
    unlock(A);  
    unlock(B).
```

The Two-Phase Locking Protocol: Example

| T_5 | T_6 | T_7 |
|---|---|------------------------------|
| lock-X(A) read(A) lock-S(B) read(B) write(A) unlock(A) | lock-X(A) read(A) write(A) unlock(A) | lock-S(A) read(A) |

The Two-Phase Locking Protocol: Exercise

18.2 Consider the following two transactions:

```
 $T_{34}$ : read( $A$ );  
      read( $B$ );  
      if  $A = 0$  then  $B := B + 1$ ;  
      write( $B$ ).
```

```
 $T_{35}$ : read( $B$ );  
      read( $A$ );  
      if  $B = 0$  then  $A := A + 1$ ;  
      write( $A$ ).
```

Add lock and unlock instructions to transactions T_{31} and T_{32} so that they observe the two-phase locking protocol.

T_{34} :

```
lock-S( $A$ )  
read( $A$ )  
lock-X( $B$ )  
read( $B$ )  
if  $A = 0$  then  $B := B + 1$   
write( $B$ )  
unlock( $A$ )  
unlock( $B$ )
```

T_{35} :

```
lock-S( $B$ )  
read( $B$ )  
lock-X( $A$ )  
read( $A$ )  
if  $B = 0$  then  $A := A + 1$   
write( $A$ )  
unlock( $A$ )  
unlock( $B$ )
```

The Two-Phase Locking Protocol: Exercise

| T1(Without locks) | T2 (Without Locks) |
|-------------------|--------------------|
| Read(A) | Read(A) |
| Read(B) | Read(S) |
| Read(C) | Read(T) |
| $A=B+C$ | $A=S+T$ |
| Write(A) | Write(A) |

The Two-Phase Locking Protocol: Exercise

| T1(Without locks) | T2 (Without Locks) |
|-------------------|--------------------|
| Read(A) | Read(A) |
| Read(B) | Read(S) |
| Read(C) | Read(T) |
| $A=B+C$ | $A=S+T$ |
| Write(A) | Write(A) |

| T1(With Locks) | T2(With Locks) |
|----------------|----------------|
| Lock-X(A) | Lock-X(A) |
| Read(A) | Read(A) |
| Lock-S(B) | Lock-S(S) |
| Read(B) | Read(S) |
| Lock-S(C) | Lock-S(T) |
| Read(C) | Read(T) |
| $A=B+C$ | $A=S+T$ |
| Write(A) | Write(A) |
| Unlock-X(A) | Unlock-X(A) |
| Unlock-S(B) | Unlock-S(S) |
| Unlock-S(C) | Unlock-S(T) |